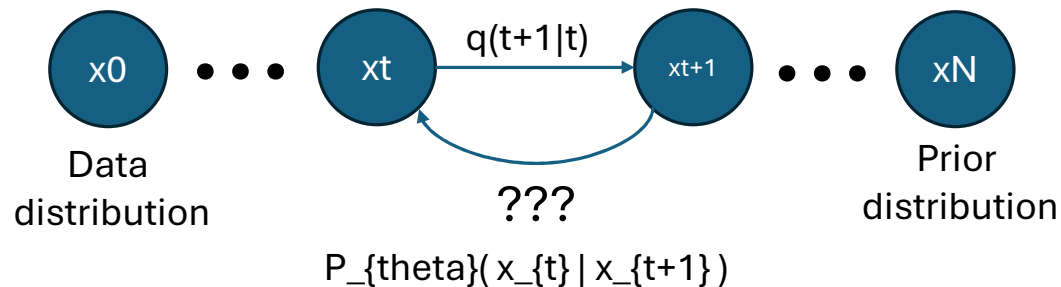


# Discretized diffusion process

- DDPM presented a discretized version of the diffusion process which presented in a Markovian framework. Diffusion occurs in steps with the following state transition kernel.

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) := \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$



- The optimization problem for the reverse problem is given by the maximizing the log-likelihood of the reversed process for all  $k$  in  $0, \dots, T$ . It can be proven, further evaluation gives the “Denoising score matching loss”.

$$\nabla \log p_k(x_k) = \mathbb{E}_{p_{0|k}(\cdot|x_k)}[\nabla \log p_{k|0}(x_k|X_0)]$$

$$Y = \mathbb{E}[X|U] \text{ if } Y = f(U), \text{ with } f = \arg \min \{\mathbb{E}[\|X - f(U)\|^2] : f \in L^2(U)\}.$$

$$\nabla \log p_k = \arg \min \{\mathbb{E}[\|f(X_k) - \nabla \log p_{k|0}(X_k|X_0)\|^2] : f \in L^2(p_k)\} .$$

# Reverse diffusion

*Denoising Diffusion Probabilistic Model*(2019)(DDPM) introduced Langevin Dynamics to solve the reverse diffusion process based on the assumption that it is a markov process. Forward models of diffusion is given as,

$$x_t = \sqrt{1 - \beta_t}x_{t-1} + \beta_t\epsilon, \quad \epsilon \sim \mathcal{N}(0, \mathbf{I})$$

Reverse sampling is called ancestral sample and closely resembles Langevin dynamics

$$x_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(x_t, t) \right) + \sigma_t \epsilon, \quad \epsilon \sim \mathcal{N}(0, \mathbf{I})$$

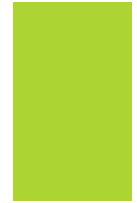
# Reverse diffusion continued...

*Denoising Diffusion Implicit Model*(2022)(DDIM) used the same approach for training the model but rejects the Markovian approach to sample and samples only a subsequence  $\tau = [\tau_{t_1} \dots \tau_{t_N}]$  with  $|\tau|$  selected as a hyper-parameter. Forward models of this diffusion process is same as that of DDPM, reverse sampling is given by,

$$x_{\tau_{i-1}} = \frac{\sqrt{\alpha_{\tau_{i-1}}}}{\sqrt{\alpha_{\tau_i}}} (x_{\tau_i} - \sqrt{1 - \bar{\alpha}_{\tau_i}} \epsilon_{\theta}(x_{\tau_i}, \tau_i)) + \sqrt{1 - \alpha_{\tau_i} - \sigma_{\tau_i}(\eta)^2} \epsilon_{\theta}(x_{\tau_i}, \tau_i) + \sigma_{\tau_i}(\eta) \epsilon$$

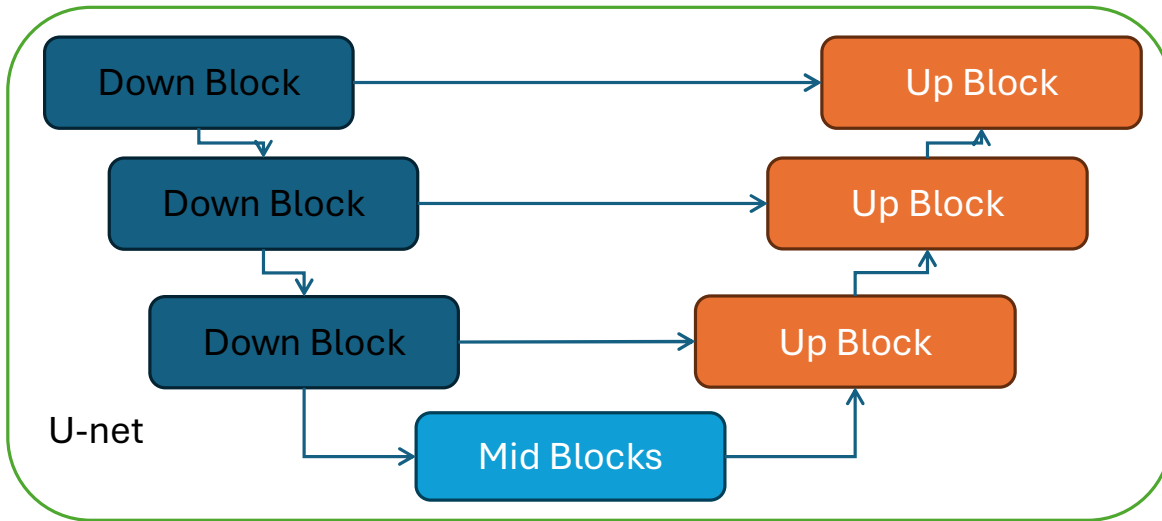
$\epsilon \sim \mathcal{N}(0, \mathbf{I})$  and,

$$\sigma_{\tau_i}(\eta) = \eta \sqrt{\frac{1 - \alpha_{\tau_{i-1}}}{1 - \alpha_{\tau_i}}} \left( \sqrt{1 - \frac{\alpha_{\tau_i}}{\alpha_{\tau_{i-1}}}} \right)$$



## Deep diffusion model using pytorch

Neural network used:



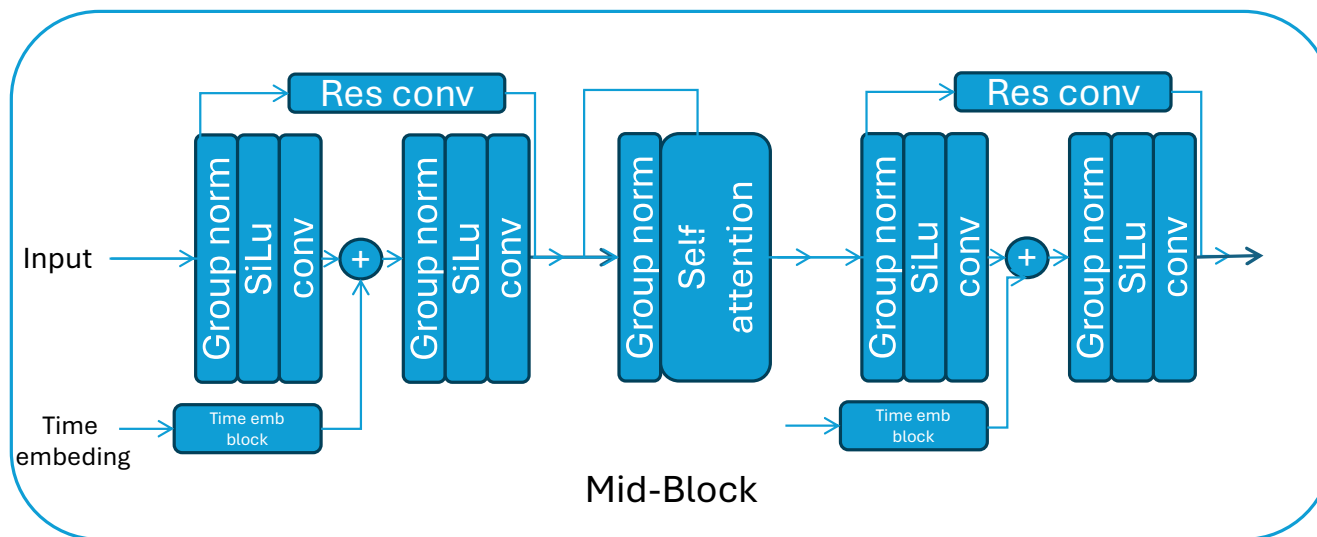
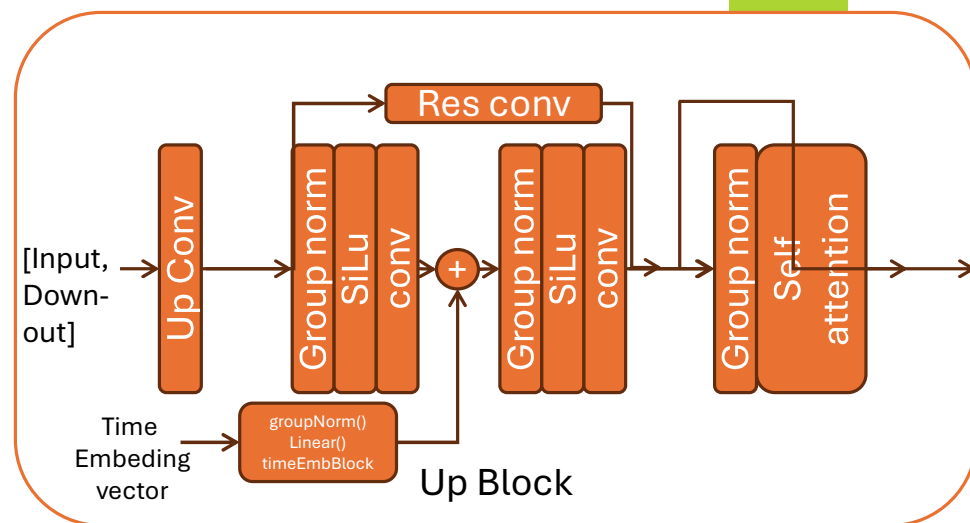
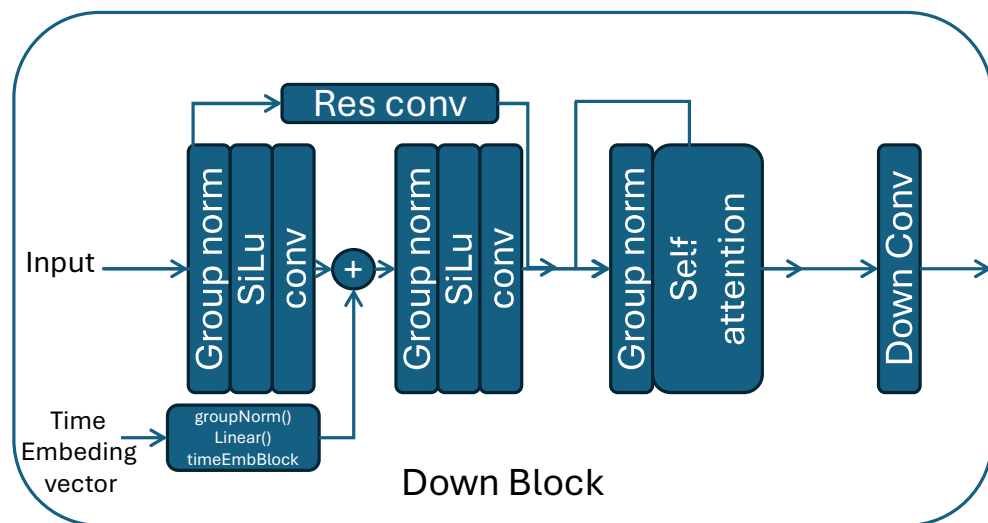
Train step:

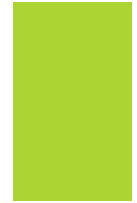
```
optimizer.zero_grad()
im = im.float().to(device)

# Sample random noise
noise = torch.randn_like(im).to(device)

# Sample timestep
t = torch.randint(0, diffusion_config['num_timesteps'], (im.shape[0],)).to(device)
noisy_im = scheduler.add_noise(im, noise, t)
noise_pred = model(noisy_im, t)

loss = criteria(noise_pred, noise)
losses.append(loss.item())
loss.backward()
optimizer.step()
```





## Deep diffusion model using pytorch

### Noise Scheduler: Add Noise(DDPM)

```
def add_noise(self, original, noise, t):
    original_shape = original.shape
    batch_size = original_shape[0]

    sqrt_alpha_cum_prod = self.sqrt_alpha_cum_prod[t]
    sqrt_one_minus_alpha_cum_prod = self.sqrt_one_minus_alpha_cum_prod[t]

    for _ in range(len(original_shape)-1):
        sqrt_alpha_cum_prod = sqrt_alpha_cum_prod.unsqueeze(-1)
        sqrt_one_minus_alpha_cum_prod = sqrt_one_minus_alpha_cum_prod.unsqueeze(-1)

    return sqrt_alpha_cum_prod*original + sqrt_one_minus_alpha_cum_prod*noise
```

### Noise Scheduler: Sample previous step(DDPM)

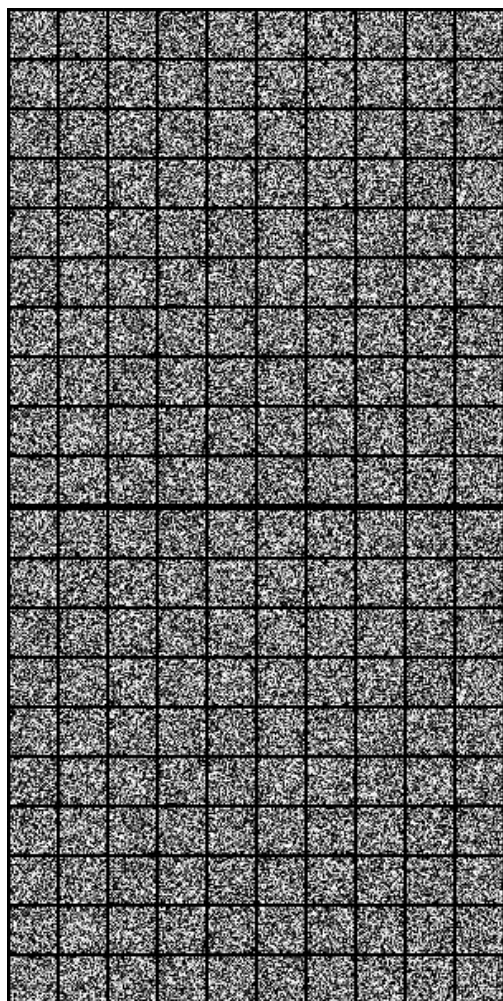
```
def sample_prev_timestep(self, xt, noise_pred, t):
    x0 = (xt - (self.sqrt_one_minus_alpha_cum_prod[t] * noise_pred)) / self.sqrt_alpha_cum_prod[t]
    x0 = torch.clamp(x0, -1, 1)

    mean = xt - ((self.betas[t] * noise_pred) / (self.sqrt_one_minus_alpha_cum_prod[t]))
    mean = mean / self.sqrt_alpha_cum_prod[t]

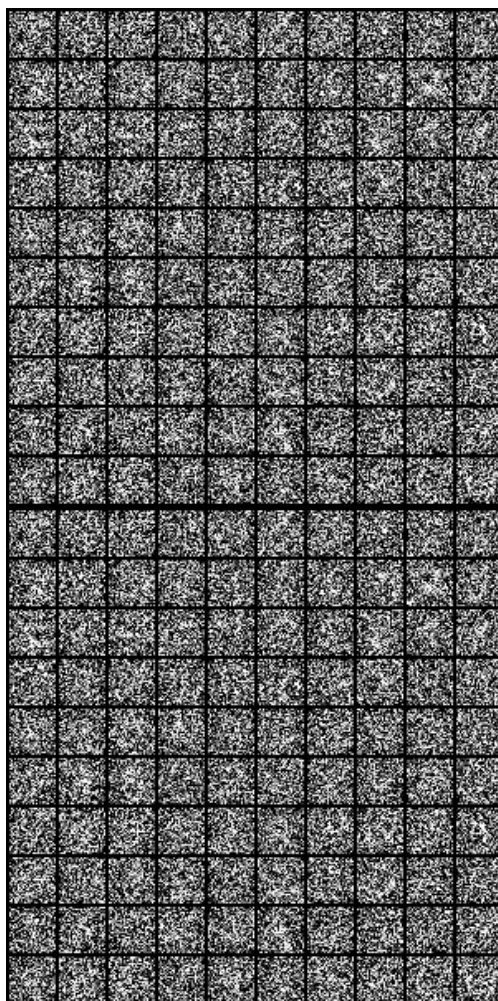
    if t==0:
        return mean, x0
    else:
        variance = ((self.betas[t])*(1 - self.alpha_cum_prod[t-1])) / (1 - self.alpha_cum_prod[t])
        sigma = variance ** 0.5
        z = torch.randn_like(xt).to(xt.device)
        return mean + sigma * z, x0
```



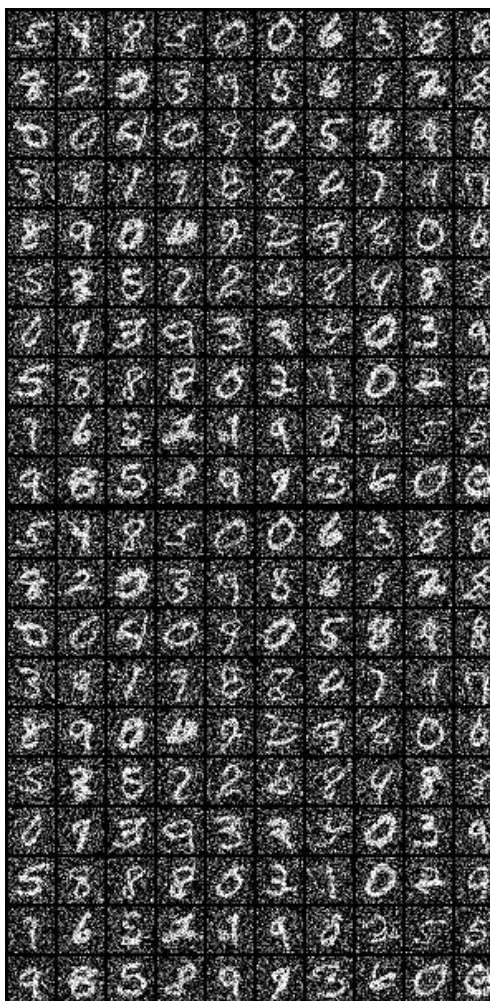
Results:



Step 750



Step 500



Step 250



Step 0